

Theano_Second

March 7, 2016

```
In [13]: import numpy as np
In [84]: import theano
import theano.tensor as T

k = T.iscalar("k")
A = T.vector("A")

# Symbolic description of the result
result, updates = theano.scan(fn=lambda prior_result, A: prior_result * A,
                              outputs_info=T.ones_like(A),
                              non_sequences=A,
                              n_steps=k)

# We only care about A**k, but scan has provided us with A**1 through A**k.
# Discard the values that we don't care about. Scan is smart enough to
# notice this and not waste memory saving them.
final_result = result[-1]

# compiled function that returns A**k
power = theano.function(inputs=[A,k], outputs=final_result, updates=updates)

print "We will Come back to this and understand it in detail"
```

We will Come back to this and understand it in detail

```
In [85]: print(power(range(10),2))
print(power(range(1),4))

[ 0.  1.  4.  9. 16. 25. 36. 49. 64. 81.]
[ 0.]
```

```
In [12]: # What's the lambda thing ?
fn=lambda prior_result, A: prior_result * A
fn(4,2)
# It is basically functions
```

```
Out[12]: 8
```

0.1 Accumulation

0.1.1 In the scan operator, the first parameter of the lambda function is the accumulated variable from the last iteration

```
In [82]: x = T.scalar()

acc,upd = theano.scan(fn = lambda temp : temp + 1, outputs_info=theano.shared(np.zeros(1)),n_s
```

- Temp accumulates the result.
- outputs.info stores the initial value of the accumulator
- n_steps tells the number of iteration

```
In [83]: acc_fn = theano.function([k],acc,updates=upd)
```

- Hence, the parameter here is only the number of times to add 1

```
In [29]: print acc_fn(3)
```

```
[[ 1.]
 [ 2.]
 [ 3.]]
```

Variable Accumulator

```
In [88]: x = T.scalar()
        inc = T.scalar()
        acc,upd = theano.scan(fn = lambda temp, inc : temp + inc, outputs_info=theano.shared(np.zeros(1)))
```

- **inc** is a scalar, which is added to the accumulator in each iteration
- In a for loop, you can equate **non_sequences** to variables that is used wholly in each iteration > Consider this C for loop >
- Here , a is added in each iteration. These variables go in the non_sequences block

```
In [89]: acc_fn = theano.function([inc,k],acc,updates=upd)
```

- So, now we have two parameters **inc** and **k**

```
In [90]: acc_fn(2,3)
```

```
Out[90]: array([[ 2.],
               [ 4.],
               [ 6.]])
```

- We can have more non_sequence variables > Consider the following piece of code >

```
In [59]: s = T.scalar('s')
        t = T.scalar('t')
        def odd_inc(x,s,t):
            x = x * s + t
            return x

        acc,upd = theano.scan(fn = odd_inc, outputs_info=theano.shared(np.zeros(1)),non_sequences = [s,t])

        acc_fn = theano.function([s,k,t],acc,updates=upd)
```

```
In [60]: acc_fn(2,5,4)
```

```
Out[60]: array([[ 2.],
               [ 10.],
               [ 42.],
               [ 170.],
               [ 682.]])
```

0.2 No Accumulation

If there no accumulation of results, we can set `outputs_info` to `None`. This indicates to scan that it doesn't need to pass the prior result to `fn`.

The general order of function parameters to `fn` is:

0.2.1 Dot Product

```
In [111]: s = T.vector('s')
          t = T.vector('t')

          def vec_mul(s,t):
              x = s * t
              return x

          acc,upd = theano.scan(fn = vec_mul, outputs_info = None, sequences = [s,t])

          acc_fn = theano.function([s,t],acc,updates=upd)
```

0.2.2 And then sum up the results

We do an elementwise multiplication and then add them up

Equate this with the following block of code >

```
In [115]: acc_fn(np.asarray([3,1,2]),np.asarray([2,2,2])).sum()

Out[115]: 12.0
```

0.2.3 Note the order of arguments in `vec_mul()`.

Try and see what happens if you put `s` at the beginning

```
In [121]: s = T.vector('s')
          t = T.vector('t')
          def vec_mul(s,t,x):
              v = s * t
              return v
          outputs_info = T.as_tensor_variable(np.asarray(0, np.float64))
          acc,upd = theano.scan(fn = vec_mul, outputs_info = outputs_info, sequences = [s,t])

          acc_fn = theano.function([s,t],acc,updates=upd)

In [124]: acc_fn([3,2,3],[2, 4,5])

Out[124]: array([ 6.,  8., 15.] )
```

1 Logistic Regression

```
In [129]: class LogisticRegression(object):
          """Multi-class Logistic Regression Class

          The logistic regression is fully described by a weight matrix :math:'W'
          and bias vector :math:'b'. Classification is done by projecting data
```

```

points onto a set of hyperplanes, the distance to which is used to
determine a class membership probability.
"""

def __init__(self, input, n_in, n_out):
    """ Initialize the parameters of the logistic regression

    :type input: theano.tensor.TensorType
    :param input: symbolic variable that describes the input of the
                  architecture (one minibatch)

    :type n_in: int
    :param n_in: number of input units, the dimension of the space in
                  which the datapoints lie

    :type n_out: int
    :param n_out: number of output units, the dimension of the space in
                  which the labels lie

    """
    # start-snippet-1
    # initialize with 0 the weights W as a matrix of shape (n_in, n_out)
    self.W = theano.shared(
        value=np.zeros(
            (n_in, n_out),
            dtype=theano.config.floatX
        ),
        name='W',
        borrow=True
    )
    # initialize the biases b as a vector of n_out 0s
    self.b = theano.shared(
        value=np.zeros(
            (n_out,),
            dtype=theano.config.floatX
        ),
        name='b',
        borrow=True
    )

    # symbolic expression for computing the matrix of class-membership
    # probabilities
    # Where:
    # W is a matrix where column-k represent the separation hyperplane for
    # class-k
    # x is a matrix where row-j represents input training sample-j
    # b is a vector where element-k represent the free parameter of
    # hyperplane-k
    self.p_y_given_x = T.nnet.softmax(T.dot(input, self.W) + self.b)

    # symbolic description of how to compute prediction as class whose
    # probability is maximal
    self.y_pred = T.argmax(self.p_y_given_x, axis=1)
    # end-snippet-1

```

```

# parameters of the model
self.params = [self.W, self.b]

# keep track of model input
self.input = input

def negative_log_likelihood(self, y):
    """Return the mean of the negative log-likelihood of the prediction
    of this model under a given target distribution.

    .. math::

        \frac{1}{|\mathcal{D}|} \mathcal{L}(\theta=\{W,b\}, \mathcal{D}) =
        \frac{1}{|\mathcal{D}|} \sum_{i=0}^{|\mathcal{D}|}
        \log(P(Y=y^{(i)}|x^{(i)}, W,b)) \quad ||
        \ell(\theta=\{W,b\}, \mathcal{D})

    :type y: theano.tensor.TensorType
    :param y: corresponds to a vector that gives for each example the
        correct label

    Note: we use the mean instead of the sum so that
        the learning rate is less dependent on the batch size
    """
    # start-snippet-2
    # y.shape[0] is (symbolically) the number of rows in y, i.e.,
    # number of examples (call it n) in the minibatch
    # T.arange(y.shape[0]) is a symbolic vector which will contain
    # [0,1,2,... n-1] T.log(self.p_y_given_x) is a matrix of
    # Log-Probabilities (call it LP) with one row per example and
    # one column per class LP[T.arange(y.shape[0]),y] is a vector
    # v containing [LP[0,y[0]], LP[1,y[1]], LP[2,y[2]], ...,
    # LP[n-1,y[n-1]]] and T.mean(LP[T.arange(y.shape[0]),y]) is
    # the mean (across minibatch examples) of the elements in v,
    # i.e., the mean log-likelihood across the minibatch.
    return -T.mean(T.log(self.p_y_given_x)[T.arange(y.shape[0]), y])
    # end-snippet-2

def errors(self, y):
    """Return a float representing the number of errors in the minibatch
    over the total number of examples of the minibatch ; zero one
    loss over the size of the minibatch

    :type y: theano.tensor.TensorType
    :param y: corresponds to a vector that gives for each example the
        correct label
    """

    # check if y has same dimension of y_pred
    if y.ndim != self.y_pred.ndim:
        raise TypeError(
            'y should have the same shape as self.y_pred',
            ('y', y.type, 'y_pred', self.y_pred.type)

```

```

    )
    # check if y is of the correct datatype
    if y.dtype.startswith('int'):
        # the T.neq operator returns a vector of 0s and 1s, where 1
        # represents a mistake in prediction
        return T.mean(T.neq(self.y_pred, y))
    else:
        raise NotImplementedError()

```

Create an object

```

In [130]: # generate symbolic variables for input (x and y represent a
           # minibatch)
           x = T.matrix('x') # data, presented as rasterized images
           y = T.ivector('y') # labels, presented as 1D vector of [int] labels

           # construct the logistic regression class
           # Each MNIST image has size 28*28
           classifier = LogisticRegression(input=x, n_in=28 * 28, n_out=10)

```

```

In [131]: # the cost we minimize during training is the negative log likelihood of
           # the model in symbolic format
           cost = classifier.negative_log_likelihood(y)

```

1.0.1 Get the gradients for bias and weight

```

In [133]: g_W = T.grad(cost=cost, wrt=classifier.W)
           g_b = T.grad(cost=cost, wrt=classifier.b)

```

1.0.2 Specify the updates

```

In [137]: # specify how to update the parameters of the model as a list of
           # (variable, update expression) pairs.
           learning_rate = T.scalar('lr')
           updates = [(classifier.W, classifier.W - learning_rate * g_W),
                      (classifier.b, classifier.b - learning_rate * g_b)]

```

```

In [163]: from logistic_sgd import load_data
           import os

           dataset='mnist.pkl.gz'

           datasets = load_data(dataset)
           batch_size = 500

           train_set_x, train_set_y = datasets[0]
           valid_set_x, valid_set_y = datasets[1]
           test_set_x, test_set_y = datasets[2]

           # compute number of minibatches for training, validation and testing
           n_train_batches = train_set_x.get_value(borrow=True).shape[0] // batch_size
           n_valid_batches = valid_set_x.get_value(borrow=True).shape[0] // batch_size
           n_test_batches = test_set_x.get_value(borrow=True).shape[0] // batch_size

```

```

index = T.lscalar('index')
# compiling a Theano function 'train_model' that returns the cost, but in
# the same time updates the parameter of the model based on the rules
# defined in 'updates'
train_model = theano.function(
    inputs=[index, learning_rate],
    outputs=cost,
    updates=updates,
    givens={
        x: train_set_x[index * batch_size: (index + 1) * batch_size],
        y: train_set_y[index * batch_size: (index + 1) * batch_size]
    }
)

```

... loading data

1.0.3 Lets create the validation and the test function

```

In [165]: test_model = theano.function(
    inputs=[index],
    outputs=classifier.errors(y),
    givens={
        x: test_set_x[index * batch_size: (index + 1) * batch_size],
        y: test_set_y[index * batch_size: (index + 1) * batch_size]
    }
)

```

```

validate_model = theano.function(
    inputs=[index],
    outputs=classifier.errors(y),
    givens={
        x: valid_set_x[index * batch_size: (index + 1) * batch_size],
        y: valid_set_y[index * batch_size: (index + 1) * batch_size]
    }
)

```

```

In [211]: test_losses = [test_model(i)
                        for i in range(n_test_batches)]
test_score = np.mean(test_losses)

```

In [212]: test_score

Out[212]: 0.2036

```

In [172]: validation_losses = [validate_model(i)
                        for i in range(n_valid_batches)]
this_validation_loss = np.mean(validation_losses)

```

```

In [210]: for i in range(100):
    for minibatch_index in range(n_train_batches):
        minibatch_avg_cost = train_model(minibatch_index, 0.0001)
    test_losses = [test_model(i)
                for i in range(n_test_batches)]
    test_score = np.mean(test_losses)
    print test_score

```

0.2036

```
In [222]: import matplotlib.pyplot as plt
def predict(i):
    """
    An example of how to load a trained model and use it
    to predict labels.
    """

    # load the saved model
    classifier = pickle.load(open('best_model.pkl'))

    # compile a predictor function
    predict_model = theano.function(
        inputs=[classifier.input],
        outputs=classifier.y_pred)


    # We can test it on some examples from test test
    dataset='mnist.pkl.gz'
    datasets = load_data(dataset)
    test_set_x, test_set_y = datasets[2]
    test_set_x = test_set_x.get_value()
    plt.imshow(test_set_x[i-1].reshape(28,28), cmap='Greys')
    plt.show()

    predicted_values = predict_model(test_set_x[i-1:i])
    return predicted_values

In [224]: %matplotlib inline
with open('best_model.pkl', 'wb') as f:
    pickle.dump(classifier, f)

number = 103
a = predict(number)
print "The label for the figure is "+ str(a)

... loading data
```

The label for the figure is [5]

1.1 Let's create a RNN

```
In [347]: import matplotlib.pyplot as plt
import numpy as np
import theano
import theano.tensor as T
import sys

dtype = theano.config.floatX
learning_rate = 0.05
n_in = 256
n_hid = 1000
n_out = 256
n_epoch = 10

def sample_weights(SizeX, SizeY):
    values = np.ndarray([SizeX, SizeY], dtype = dtype)
    for dx in range(SizeX):
        row_val = np.random.normal(loc = 0.0, scale = 0.1, size=(SizeY,))
        values[dx,:] = row_val
    _,svs,_ = np.linalg.svd(values)
    values = values / sv[0]
    return values

In [348]: def generate_parameters(n_in, n_hid, n_out):
    b_out = theano.shared(np.zeros(n_out, dtype = dtype))
    b_hid = theano.shared(np.zeros(n_hid, dtype = dtype))
```

```

W_ih = theano.shared(sample_weights(n_in, n_hid))
W_hh = theano.shared(sample_weights(n_hid, n_hid))
W_ho = theano.shared(sample_weights(n_hid, n_out))
h_0 = theano.shared(np.zeros(n_hid, dtype = dtype))
return h_0, b_out, b_hid, W_ih, W_hh, W_ho

def logistic_function(vec):
    return 1/(1 + T.exp(-vec))
def activ_tan(vec):
    return T.tanh(vec)
def one_step(x_t, hid_s, W_ih, W_hh, W_ho, b_out, b_hid):
    h_t = activ_tan(theano.dot(x_t, W_ih) + theano.dot(hid_s, W_hh) + b_hid)
    return h_t

In [349]: hid_s, b_out, b_hid, W_ih, W_hh, W_ho = generate_parameters(n_in,n_hid,n_out)

params = [ b_out, b_hid, W_ih, W_hh, W_ho]

x_t = theano.shared(np.random.uniform(size = n_in))

inp = T.matrix(dtype = dtype)

target = T.matrix(dtype = dtype)

hidden_s,_ = theano.scan(fn = one_step, sequences=inp, outputs_info = hid_s, non_sequences =
#hid_s = hidden_s[hidden_s.shape[0]-1]

y_t = theano.dot(hidden_s[hidden_s.shape[0]-1], W_ho) + b_out

p_y_given_x = T.nnet.softmax(y_t)

y_t = T.argmax(p_y_given_x, axis = 1)

lr = theano.shared(np.cast[dtype](learning_rate))

cost = -T.sum(target*T.log(p_y_given_x) + (1.- target)*T.log(1. - p_y_given_x))

In [350]: def get_train_graph(target, inp, cost):
    grads = []
    for param in params:
        grads.append(T.grad(cost, param))
    update = []
    for param,grad in zip(params, grads):
        update.append((param, param - grad*lr))
    train_fn = theano.function(inputs = [inp,target], outputs = cost, updates = update)
    return train_fn

In [351]: def get_pred_graph(inp):
    predictions = theano.function(inputs = [inp], outputs = y_t, updates = [(hid_s, hidden_s[
    return predictions

In [352]: def convert_string(file):
    f = open(file,'r')
    text = f.read()

```

```

f.close()
inp = np.zeros([len(text), 256], dtype=dtype)
out = np.zeros([len(text), 256], dtype=dtype)
counter = 0
for char in text:
    if(counter > 0):
        inp[counter][ord(char)] = 1
        out[counter-1][ord(char)] = 1
        counter = counter + 1
    return [inp, out]

In [353]: learn_rnn_fn = get_train_graph(target, inp, cost)
pred_rnn_fn = get_pred_graph(inp)

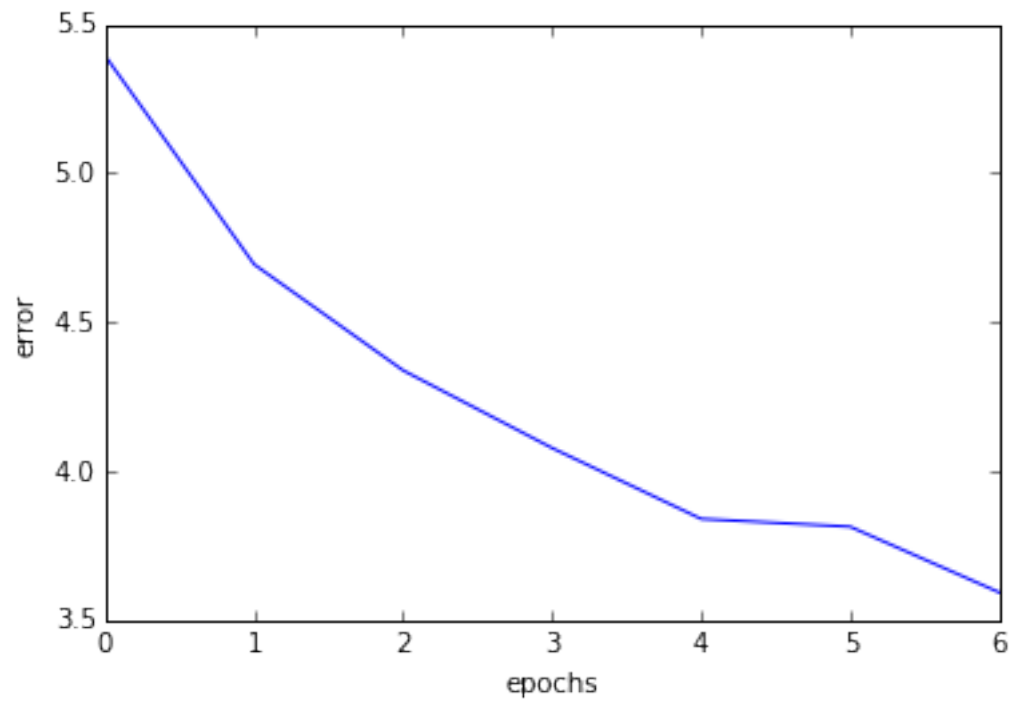
train_data = convert_string("log")

In [357]: j = 0
def train_rnn(train_data, n_epoch = 100):
    train_err = np.ndarray(n_epoch)
    for i in range(n_epoch):
        for j in range(len(train_data[0])):
            tempInp = np.zeros([1,256], dtype=dtype);
            tempInp[0] = train_data[0][j]
            tempOut = np.zeros([1,256], dtype=dtype);
            tempOut[0] = train_data[1][j]
            train_cost = learn_rnn_fn(tempInp, tempOut)
            sys.stdout.write(chr(pred_rnn_fn(tempInp)))
            train_err[i]=train_err[i]+ train_cost
            train_err[i]= train_err[i]/len(train_data[0])
        print ""
        if train_err[i] > 5*train_err[i-1]:
            break
    return train_err

In [ ]: train_errors = train_rnn(train_data, n_epoch)
for i in range(0, len(train_data[0])):
    temp = np.zeros([1,256], dtype=dtype);
    temp[0] = train_data[0][i]
    sys.stdout.write(chr(pred_rnn_fn(temp)))

In [359]: def plot_learning_curve(train_err):
    plt.plot(np.arange(n_epoch), train_errors, 'b-')
    plt.xlabel('epochs')
    plt.ylabel('error')
    plt.show()
    plot_learning_curve(train_errors)

```



In []:

In []: